


Martin Mundt, Dr. Iuliia Pliushch, Prof. Dr. Visvanathan Ramesh

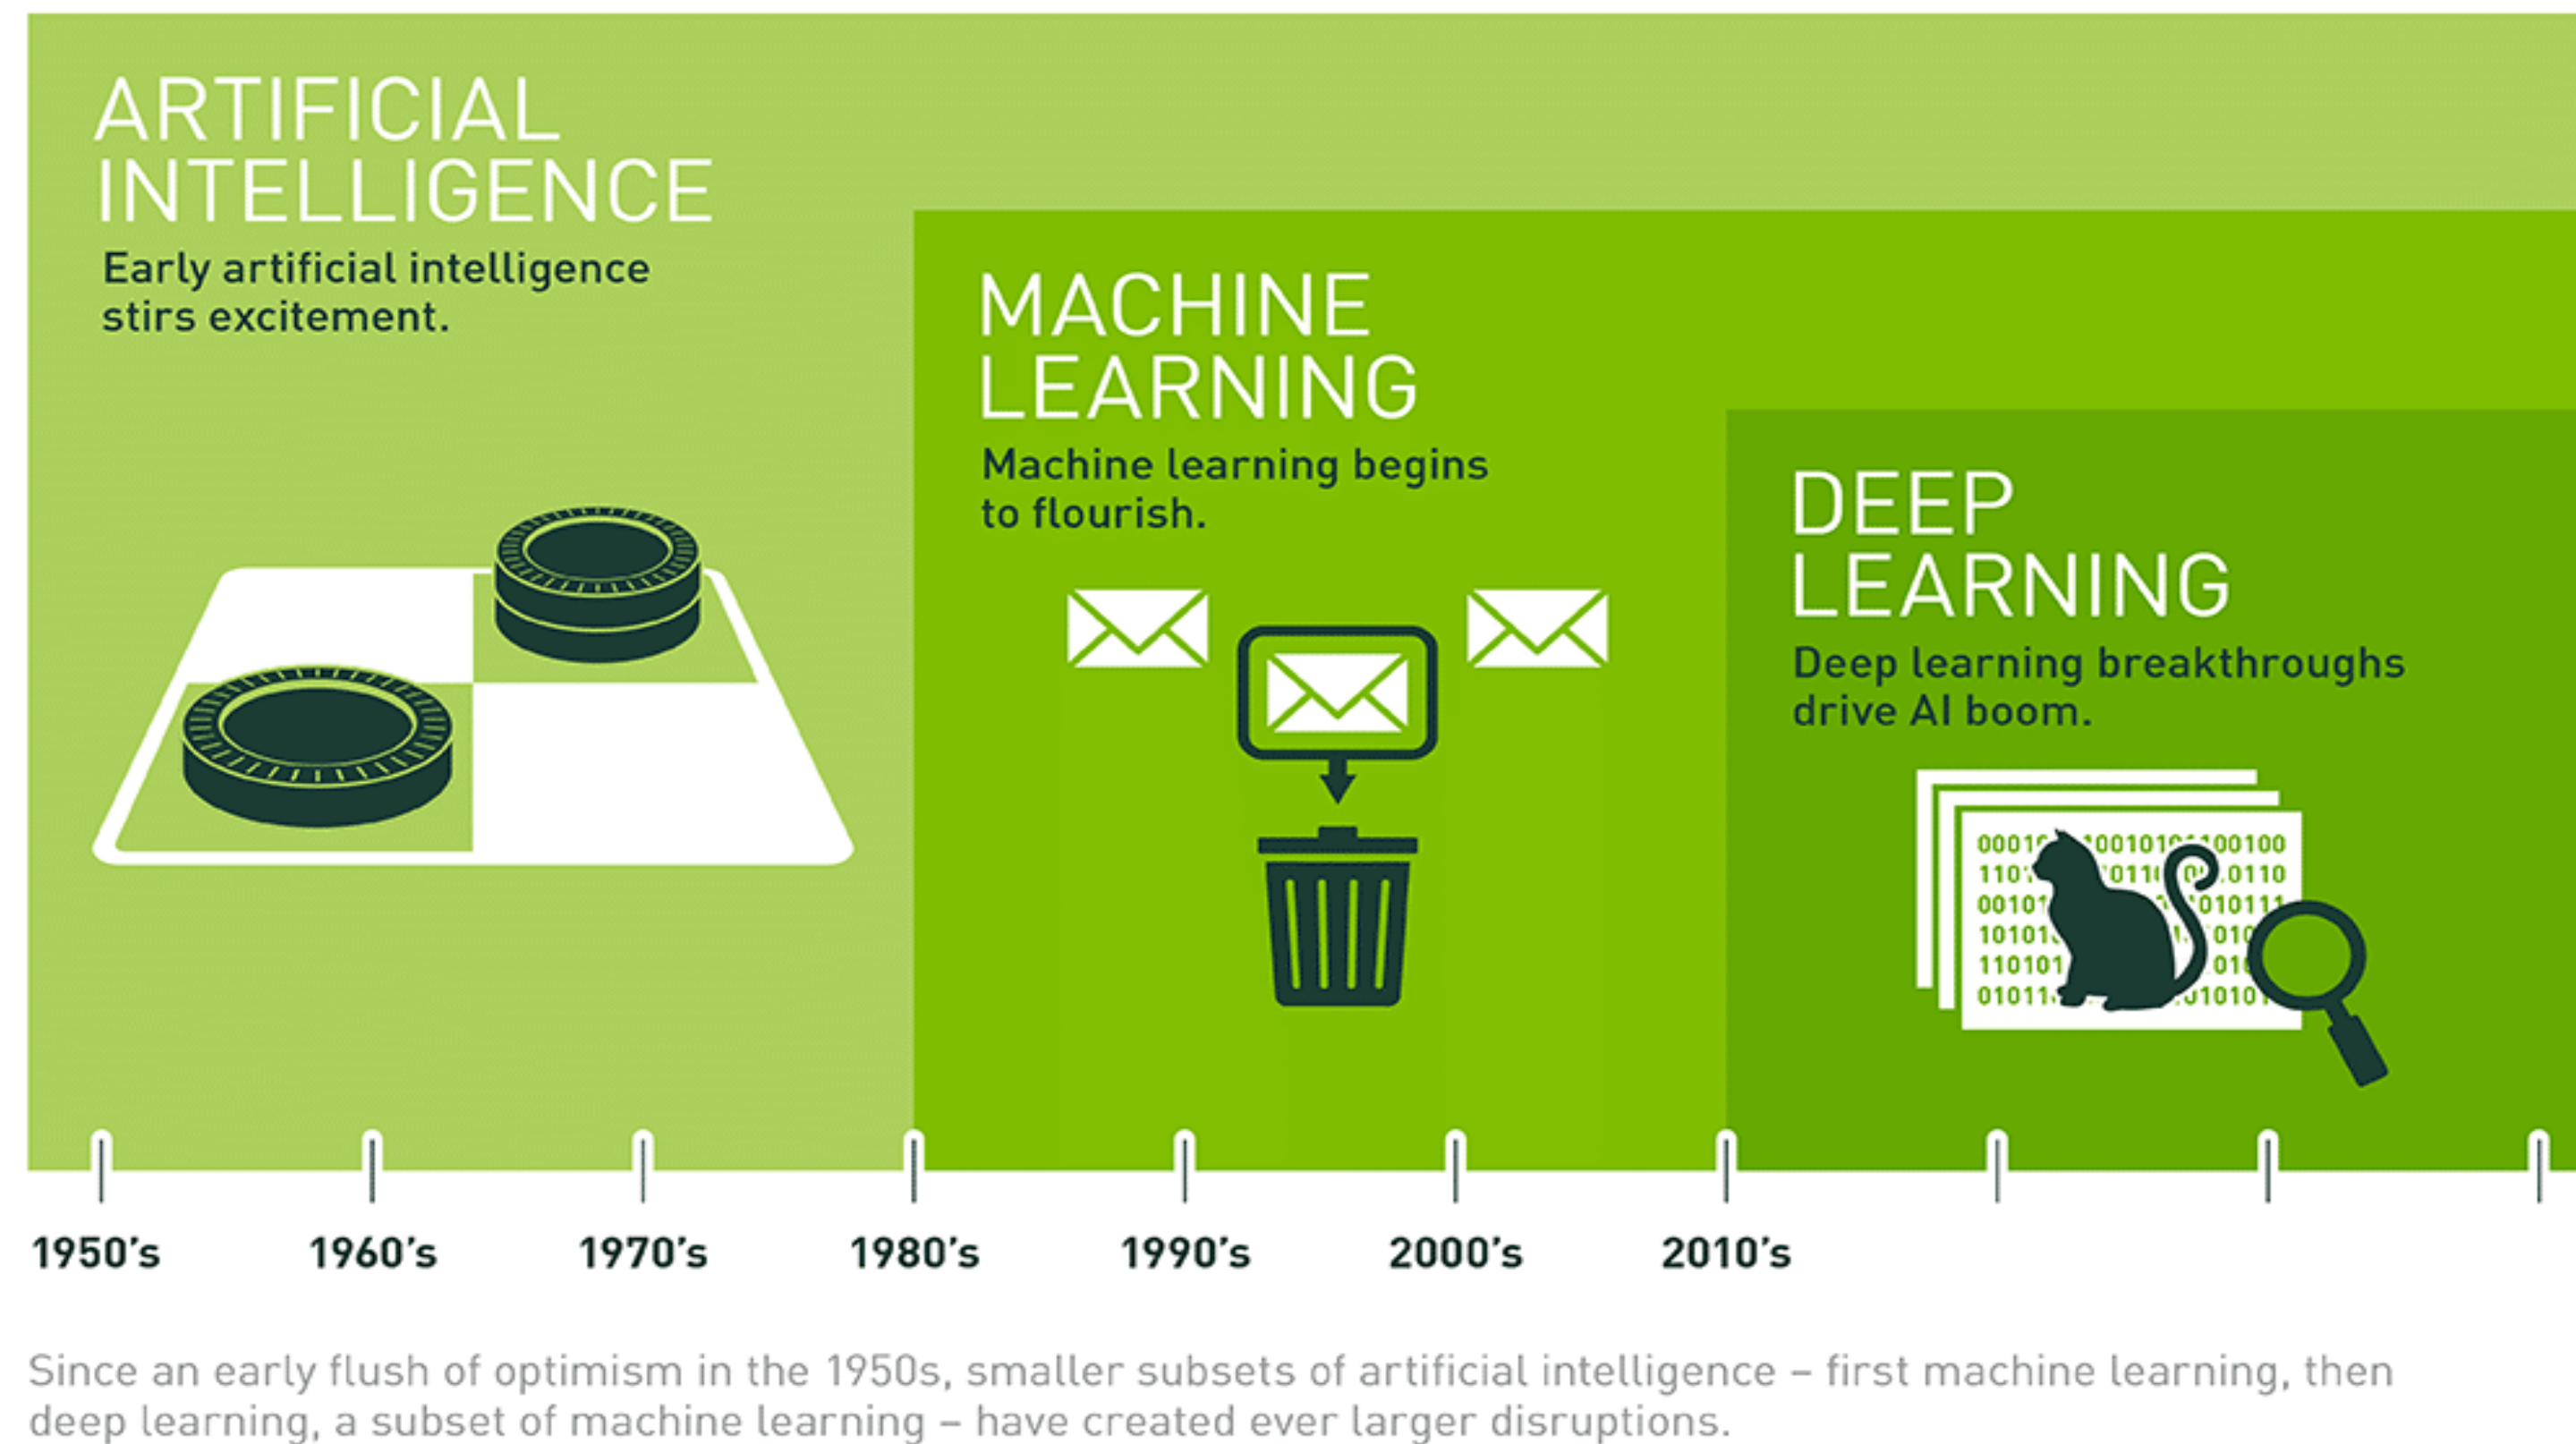
Pattern Analysis & Machine Intelligence

Praktikum: MLPR-SS21

Week 10: Introduction into Reinforcement Learning



AI, machine learning and NNs



<https://blogs.nvidia.com/blog/2018/08/02/supervised-unsupervised-learning/>

Different types of learning

TYPES OF MACHINE LEARNING



Supervised Learning

Train an algorithm to perform classification and regression with a labelled data set.



Unsupervised Learning

Train an algorithm to find clusters and associations in an unlabelled data set.




Reinforcement Learning

Train an agent to take certain actions in an environment without a data set.

<https://www.breakfreegraphics.com/design-blog/an-intro-to-machine-learning-for-Designers/>

Supervised vs. unsupervised learning




Supervised Learning

Unsupervised Learning

https://twitter.com/athena_schools/status1063013435779223553photo/1

Reinforcement learning



<https://aihub.org/2020/06/30/the-ingredients-of-real-world-robotic-reinforcement-learning/>

Reinforcement Learning: Applications


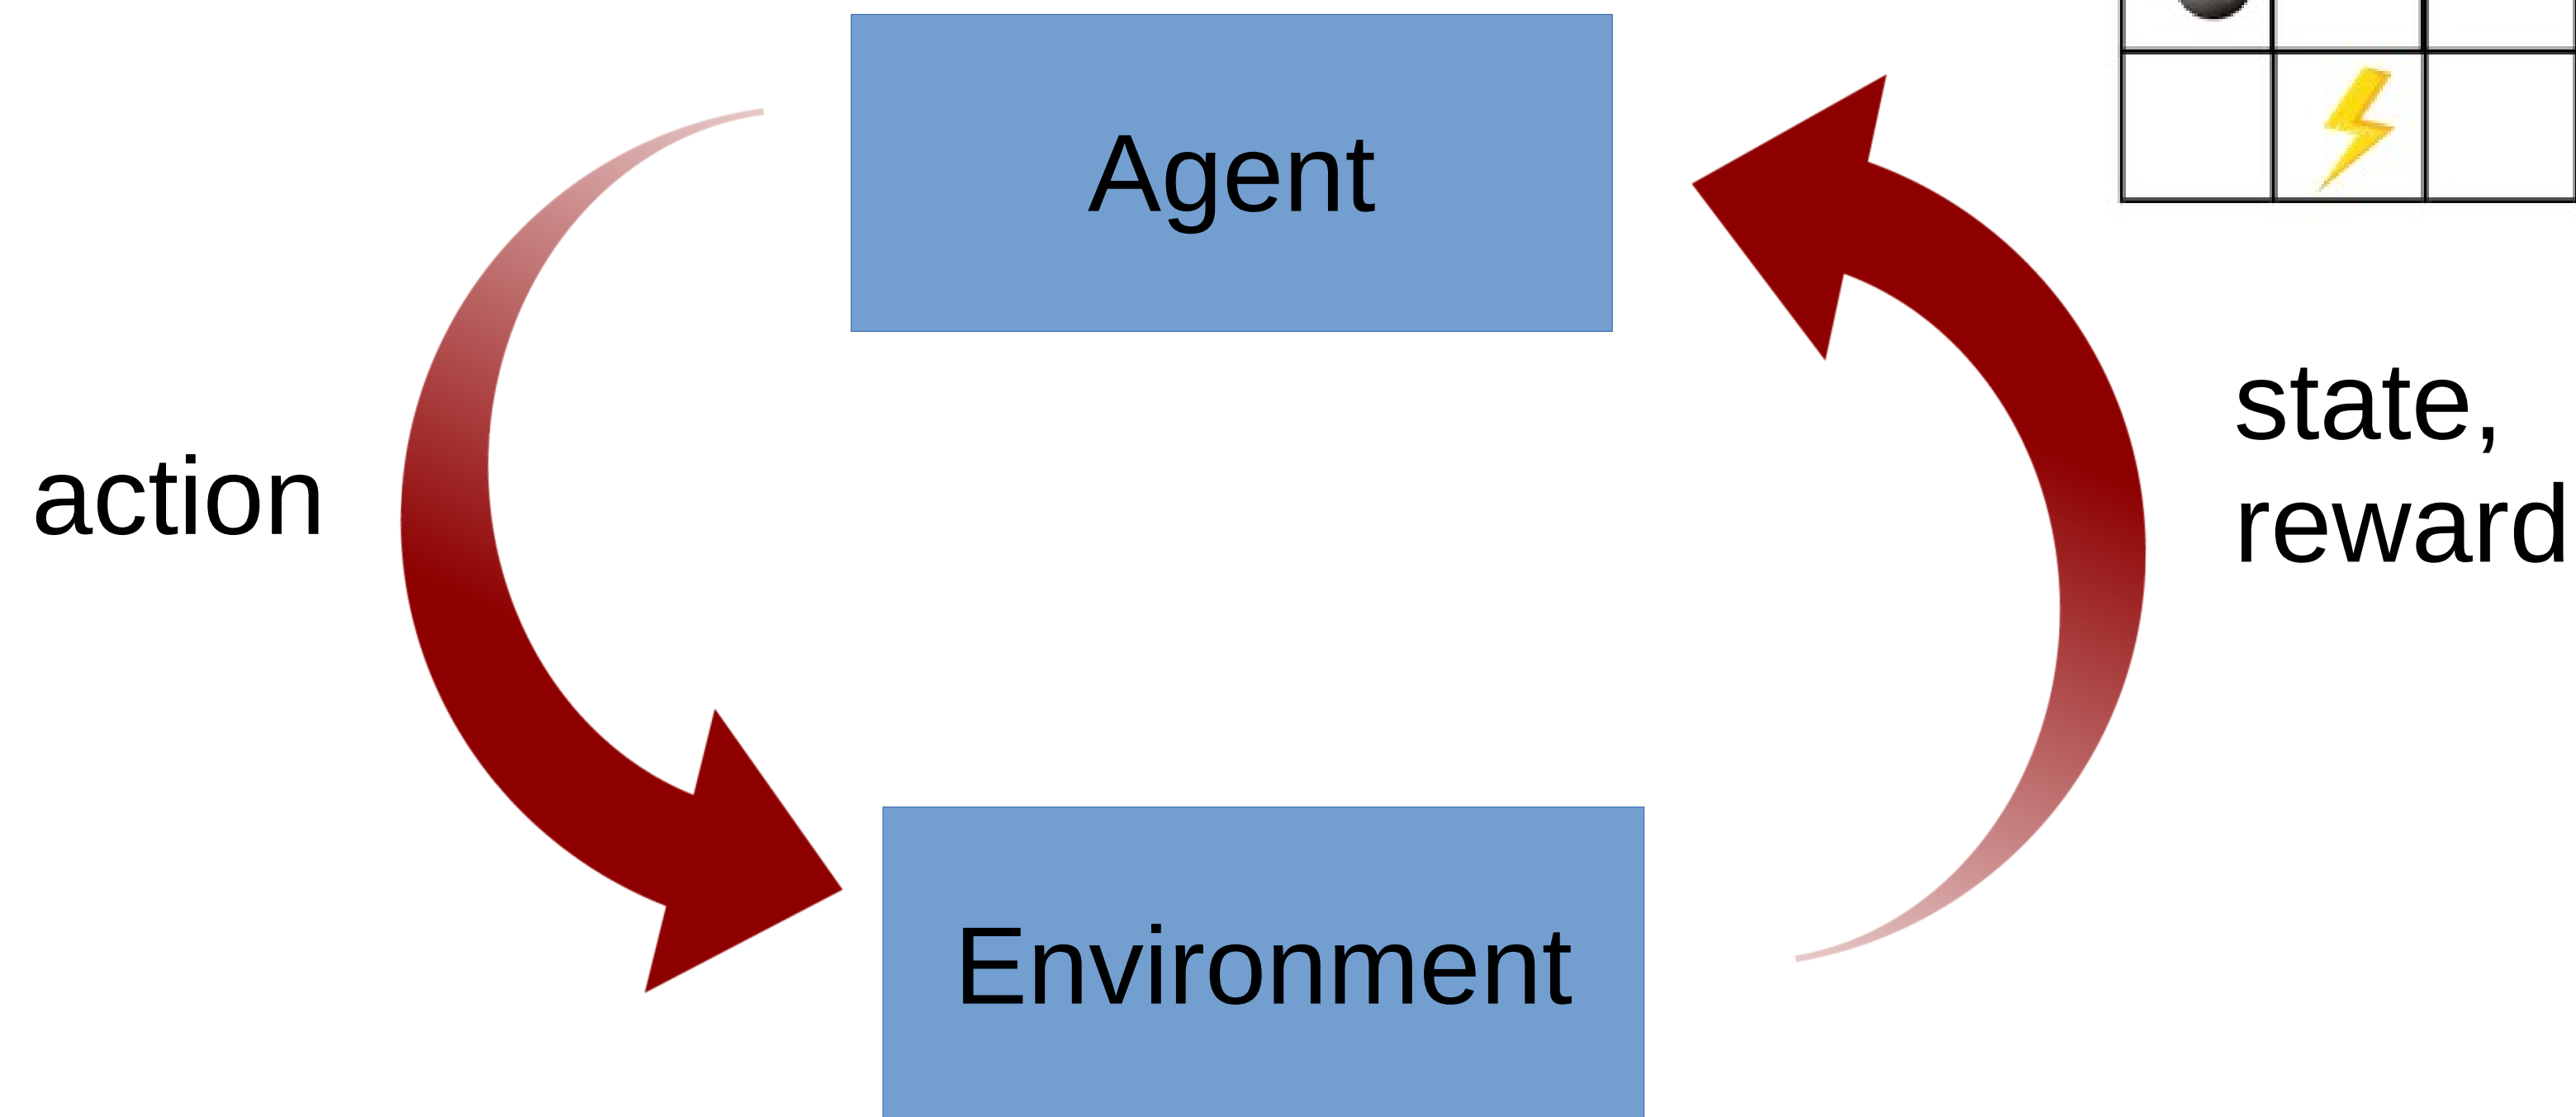
- Learning how to play games
- Robotics
- Finance
- Healthcare
- Meta-Learning



<https://ai.googleblog.com/2018/06/scalable-deep-reinforcement-learning.html>

Reinforcement Learning: Details


- Learning to maximize **rewards** by performing **actions** in an **environment**.



<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

- **Environment:** maze
- **Actions:** left/right/up/down
- **Rewards:**
 - -1 on each step,
 - -100 to step on mine
 - 1 for lightning charge
 - 100 for end
- **Policy:** mapping from states to actions, e.g.
 - always go left until wall, then right
 - after stepping on mine, always go right+down



<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

- **Q-Table:**

a table storing the expected rewards for every (state, action)-pair

Actions : ↑ → ↓ ←

Start				
Nothing / Blank				
Power				
Mines				
END				

<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

Scenario:

- a learning **agent**
- **S**: a set of possible states
- **A**: a set of possible actions
- a **state transition** function

$$\delta: S \times A \rightarrow S$$

- a **reward** function

$$r: S \times A \rightarrow \mathbb{R}$$

Feedback loop:

- the agent repeatedly chooses an action according to some **policy**

$$\pi: S \rightarrow A$$


- the environment changes to a new state according to δ
- some states provide the agent with feedback (**reinforcement**)

Reinforcement Learning: Reward

- Cumulative expected reward:

$$G_t = \sum_{i=0}^{\infty} \gamma^i * r_{t+i}$$


(γ makes the G_t finite)



- Cumulative expected reward:

$$G_t = \sum_{i=0}^{\infty} \gamma^i * r_{t+i}$$

(γ makes the G_t finite)



- Bellman **expectation** for the state-value function:


$$\begin{aligned} v_\pi(s) &= E[G_t | S_t = s] = E_\pi[R_t + \gamma * G_{t+1} | S_t = s] \\ &= \underbrace{\sum_a \pi(a|s)}_{\text{policy stochasticity}} \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * \underbrace{E_\pi[G_{t+1} | S_{t+1} = \hat{s}]}_{v_\pi(\hat{s})}] \end{aligned}$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

<https://www.ke.tu-darmstadt.de/lehre/archiv/ss09/ki/reinforcement-learning.pdf>

- State-value to action-value function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$



- Bellman **expectation** for the action-value function:

$$\begin{aligned} q_{\pi}(s, a) &= E[G_t | S_t = s, A_t = a] = E_{\pi}[R_t + \gamma * G_{t+1} | S_t = s, A_t = a] \\ &= \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * \underbrace{[r + \gamma * E_{\pi}[G_{t+1} | S_{t+1} = \hat{s}]]}_{v_{\pi}(\hat{s})} \end{aligned}$$

Reinforcement Learning: Optimal Policy

$$v_{opt}(s) = \max_{\pi} v_{\pi}(s)$$

$$\pi_{opt} = \arg \max_{\pi} v_{\pi}(s)$$

$$q_{opt}(s, a) = \max_{\pi} q_{\pi}(s, a)$$

$$\pi_{opt}(s) = \arg \max_a q_{\pi}(s, a)$$

Bellman **optimality** equations:

$$v_{opt}(s) = \max_a \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * v_{opt}(\hat{s})]$$

$$q_{opt}(s, a) = \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * \max_{\hat{a}} q_{opt}(\hat{s}, \hat{a})]$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

Reinforcement Learning: Q-Learning

Model-free (train on trajectories),
Off-policy (not train on own policy)


$$\forall s \in S, \forall a \in A, q(s, a) = 0$$

Loop:

Sample $\langle s, a, r, \hat{s} \rangle$


Compute $\tilde{q}(s, a) = r(s, a) + \gamma * \max_{a_i} q(\hat{s}, a_i)$

Update $q(s, a) = \alpha * \tilde{q}(s, a) + (1 - \alpha) * q(s, a)$



<http://icaps18.icaps-conference.org/fileadmin/alg/conferences/icaps18/summerschool/lectures/Lecture5-rl-intro.pdf>
<https://www.coursera.org/learn/practical-rl/home/welcome>

Reinforcement Learning: Q-Learning



How to sample \hat{s} ?


ϵ -greedy policy

Exploration-exploitation trade-off:

With probability ϵ choose a **random** action,
else the **best** one.

<https://www.coursera.org/learn/practical-rl/home/welcome>

Reinforcement Learning: Openai Gym



<https://gym.openai.com/envs/Taxi-v2/>

<https://www.coursera.org/learn/practical-rl/home/welcome>