

Martin Mundt, Dr. Iuliia Pliushch, Prof. Dr. Visvanathan Ramesh

Pattern Analysis & Machine Intelligence

Praktikum: MLPR-SS21

Week 13: Metalearning with Reinforce



On-policy vs. off-policy RL algorithms

<i>On-policy</i>	<i>Off-policy</i>
Agent can pick actions.	Agent can't pick actions.
Agent always follows his own policy.	Learning with exploration, playing without. For example, learning from experts or recorded sessions.
Example: reinforce (policy gradient method)	Example: Q-learning (value-based method)

<https://www.coursera.org/learn/practical-rl/home/welcome>

Policy

π – policy

s – state

a – action

Deterministic

$$a = \pi_{\theta}(a|s)$$

Stochastic

$$a \sim \pi_{\theta}(a|s)$$

Which kind of policy would you choose if you wanted to

Learn to win rock-paper-scissors?

<https://www.coursera.org/learn/practical-rl/home/welcome>

Policy gradient methods: objective

- Maximize expected reward J
- J is the expectation over states and actions of earned reward R
- $p(s)$ – probability of visiting a certain state
- $\pi_{\theta}(a|s)$ – policy stochasticity: probability of choosing an action a in state s

$$J = E_{s \sim p(s), a \sim \pi_{\theta}(a|s)} R(s, a)$$

$$= \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

State visiting frequency
(policy-dependent)

Reward for 1-step session
Assumption: the session takes only 1 step

<https://www.coursera.org/learn/practical-rl/home/welcome>

Policy gradient methods: derivative

- What is the analytical **derivative of J** which can be easily approximated with **sampling**?

$$\begin{aligned} J &= E_{s \sim p(s), a \sim \pi_{\theta}(a|s)} R(s, a) \\ &= \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds \end{aligned}$$

$$\nabla \pi(z) = \pi \nabla \log \pi(z)$$

$$\nabla J = \int_s p(s) \int_a \pi_{\theta}(a|s) \nabla \log \pi_{\theta}(s|a) R(s, a) da ds$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

Policy gradient methods: reward

- How should we change **J** if our session/game takes more than one step to complete?

$$J = E_{s \sim p(s), a \sim \pi_{\theta_0}(a|s)} \log \pi_{\theta_0}(a|s) R(s, a) da ds$$

- substitute the **reward R** with **cumulative discounted reward G**

$$G_t = \sum_{i=0}^{\infty} \gamma^i * r_{t+i}$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

Policy gradient methods: Reinforce

- initialize NN weights θ_0
- Loop:
 - Sample N sessions z_i under current $\pi_{\theta}(a|s)$
 - Evaluate policy gradient

Cumulative discounted reward

$$J = E_{s \sim p(s), a \sim \pi_{\theta}(a|s)} \log \pi_{\theta}(a|s) G(s, a) da ds =$$
$$= \frac{1}{M} \sum_{z_i \in Z} \sum_{s, a \in z_i} \log \pi_{\theta}(a|s) G(s, a)$$

- Ascend

$$\theta_{i+1} \leftarrow \theta_i + \alpha \nabla J$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

RL application: Neural architecture search

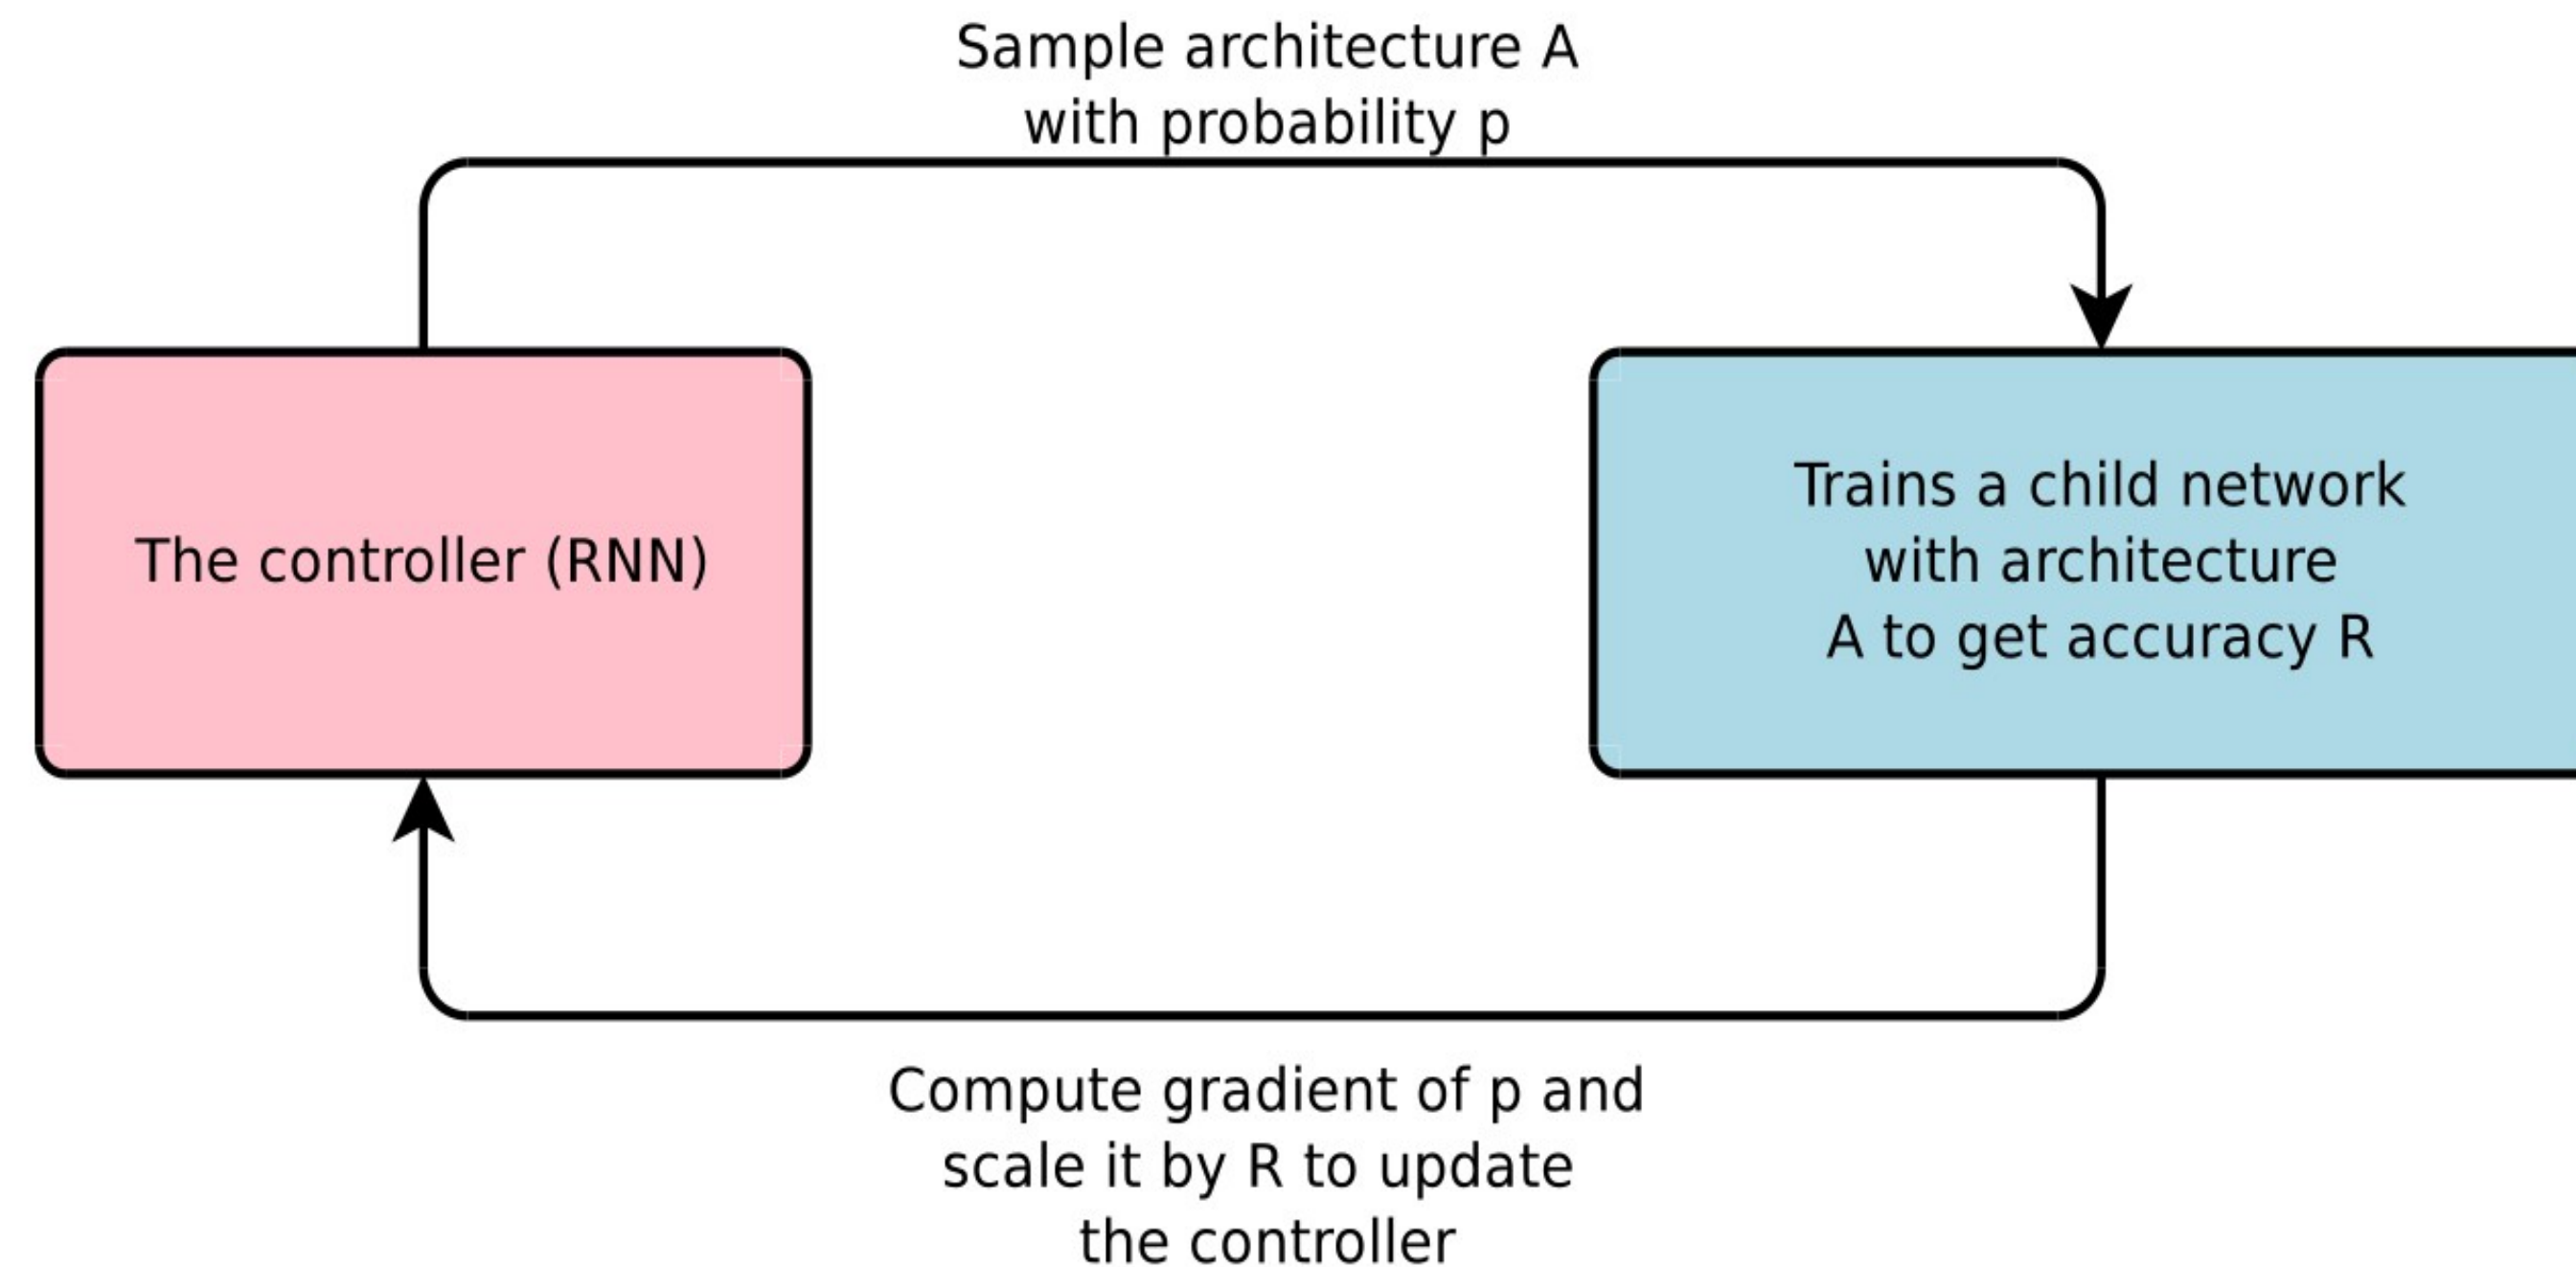
- Suppose we wanted to find out what is the best **neural architecture** to solve a problem, for example MNIST digit recognition
- How would our **environment** look? What are the **states**, **actions** and **rewards**?
- What are the neural architecture **parameters** that we could change?

States: encode parameter choices

Actions: fix parameter values

Rewards: accuracy of the network

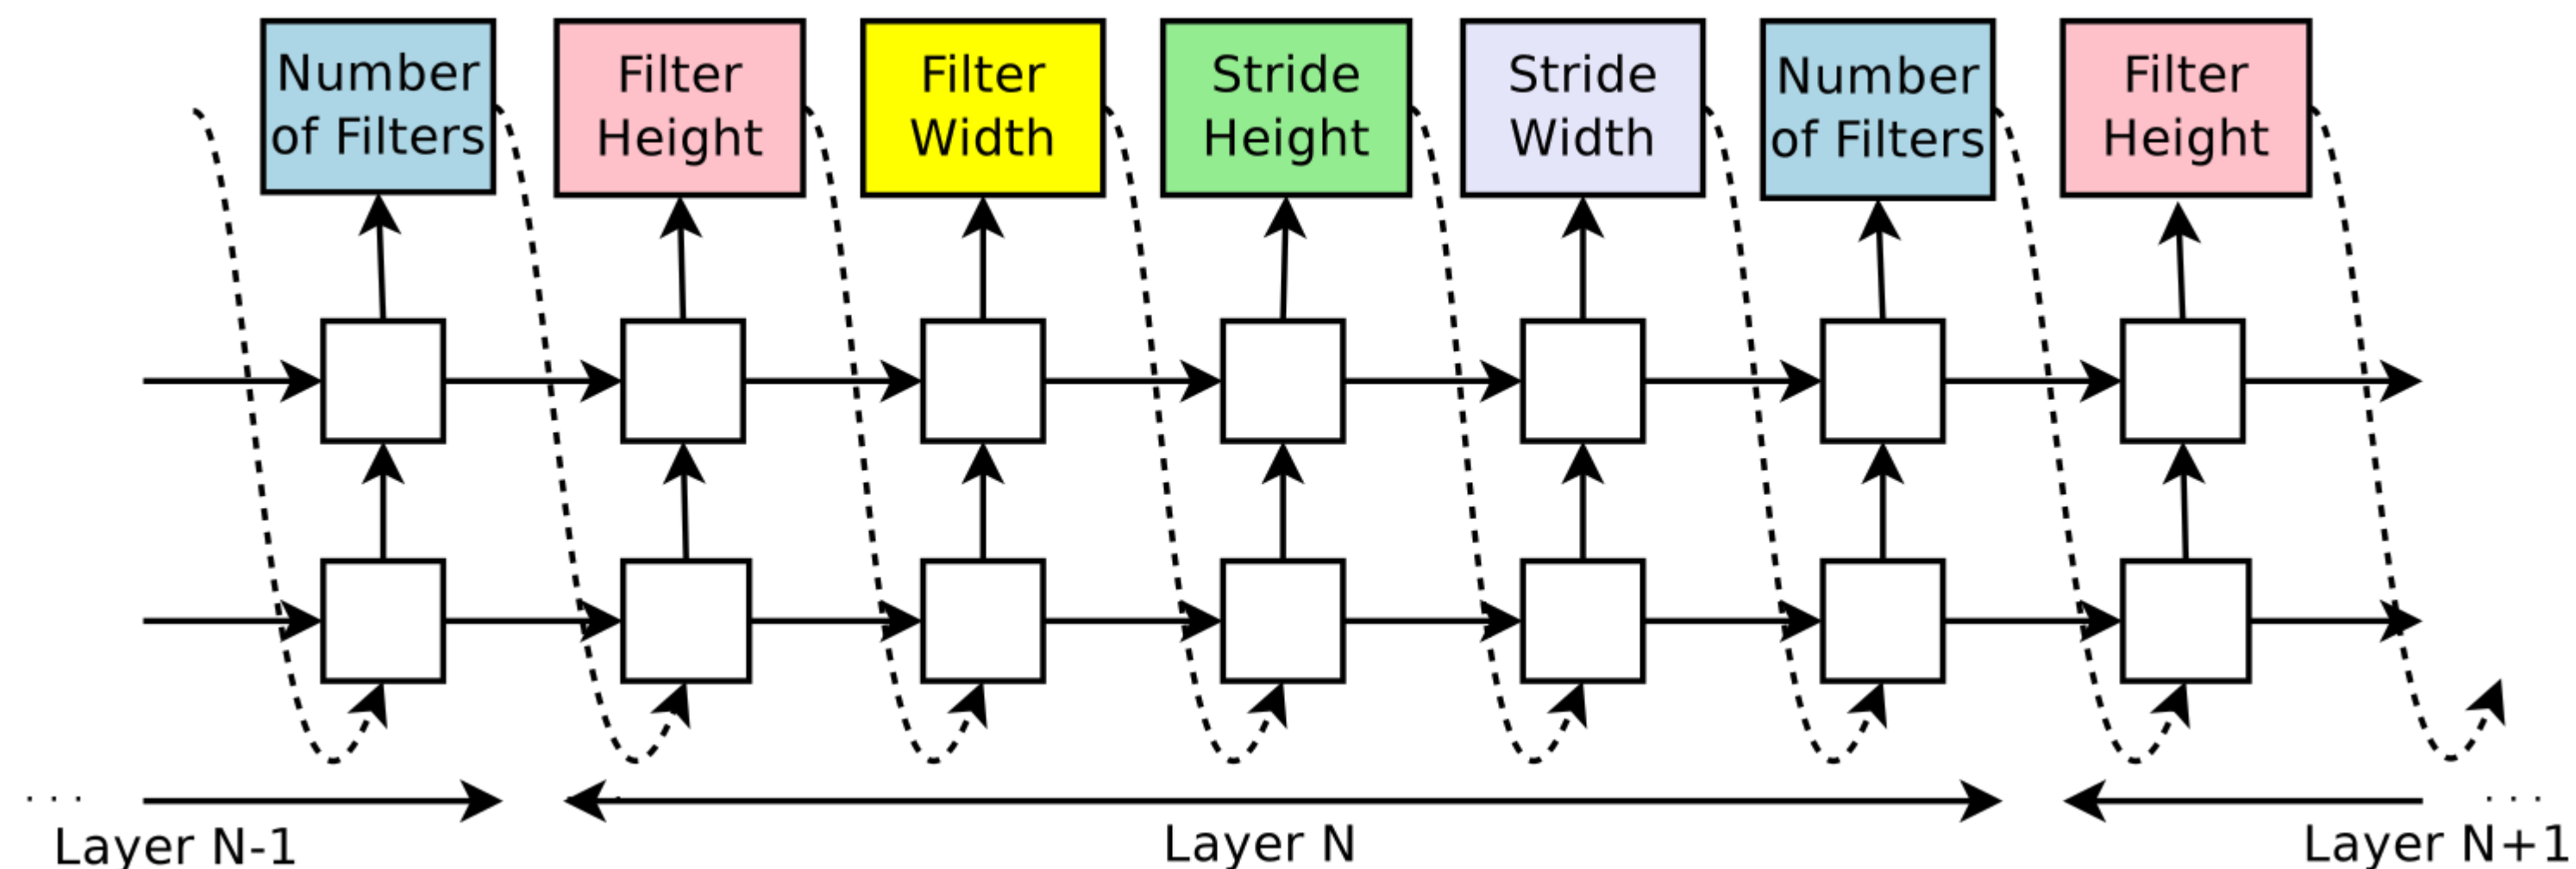
Meta-Learning: Neural architecture search



Zoph, Barret, and Quoc V. Le. "Neural architecture search with reinforcement learning." ICLR (2017).

Meta-Learning: Neural architecture search

- Suppose we use an **LSTM** cell, what is the sequence it should learn?
- The **sequence** starts with predicting the parameters of the **first layer** and ends with the **last**, such that given the sequence, a **NN** can be constructed, trained on a chosen task and the **accuracy** computed.



Zoph, Barret, and Quoc V. Le. "Neural architecture search with reinforcement learning." ICLR (2017).

Meta-Learning: Neural architecture search

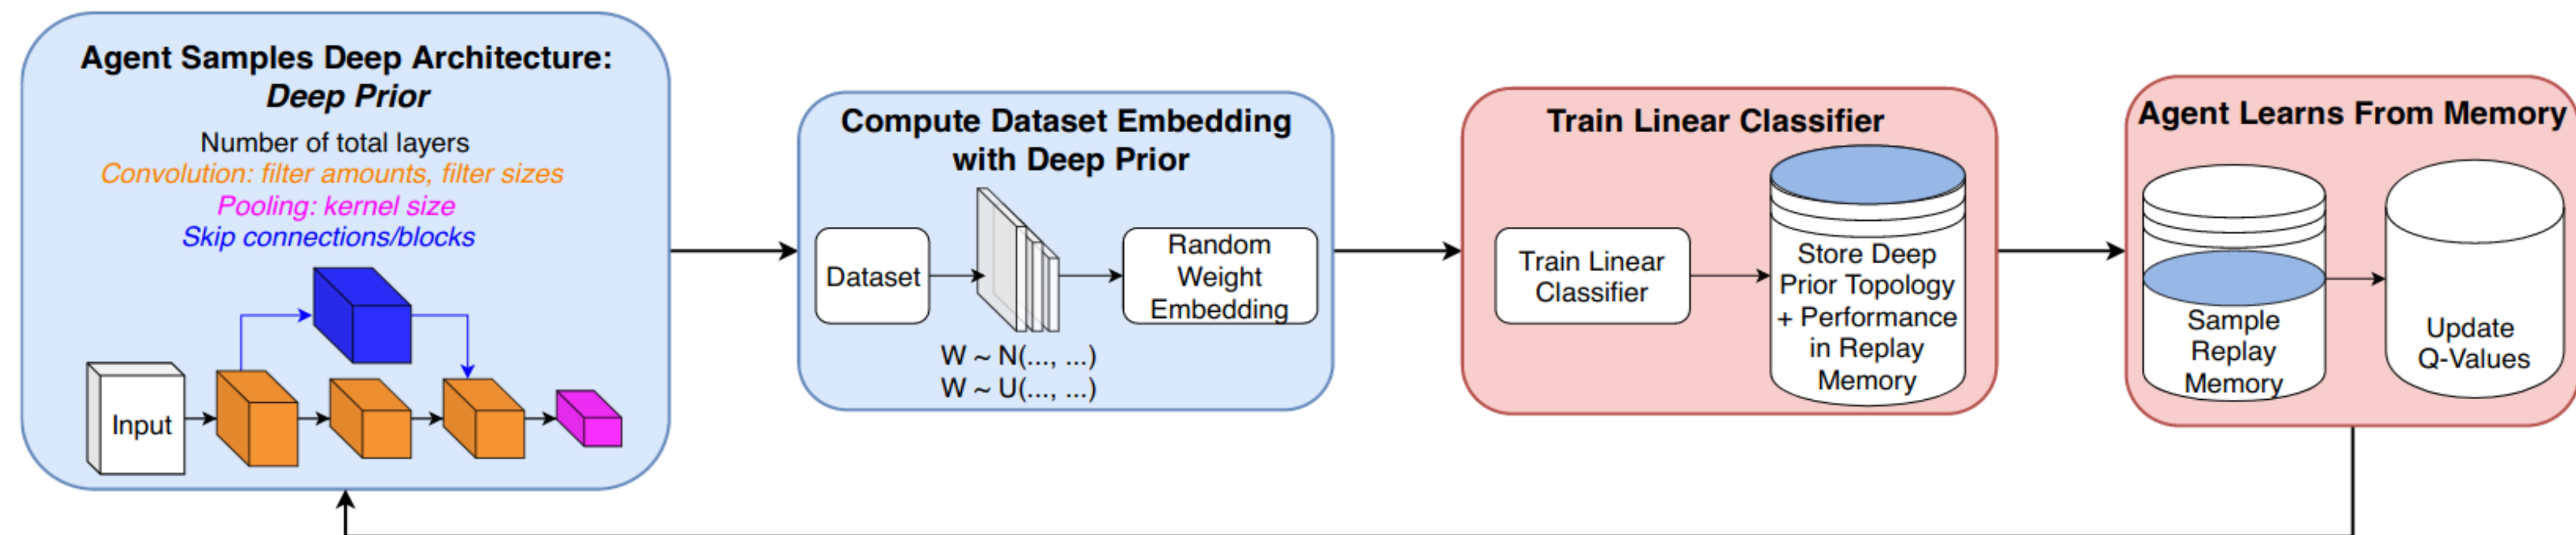
- Suppose we use an **LSTM** cell, what is the sequence it should learn?
- The **sequence** starts with predicting the parameters of the **first layer** and ends with the **last**, such that given this sequence, a **NN** can be constructed, trained on a chosen task and the **accuracy** computed.
- Hence, how often do we get **rewards** from the environment, if a step in it involves fixing **one** of the neural architecture parameters?
 - The rewards we get are **sparse**.

Meta-Learning Application: Defect classification with CODEBRIM



Mundt, M., Majumder, S., Murali, S., Panetsos, P., & Ramesh, V. (2019). Meta-learning convolutional neural architectures for multi-target concrete defect classification with the concrete defect bridge image dataset. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (pp. 11196-11205).

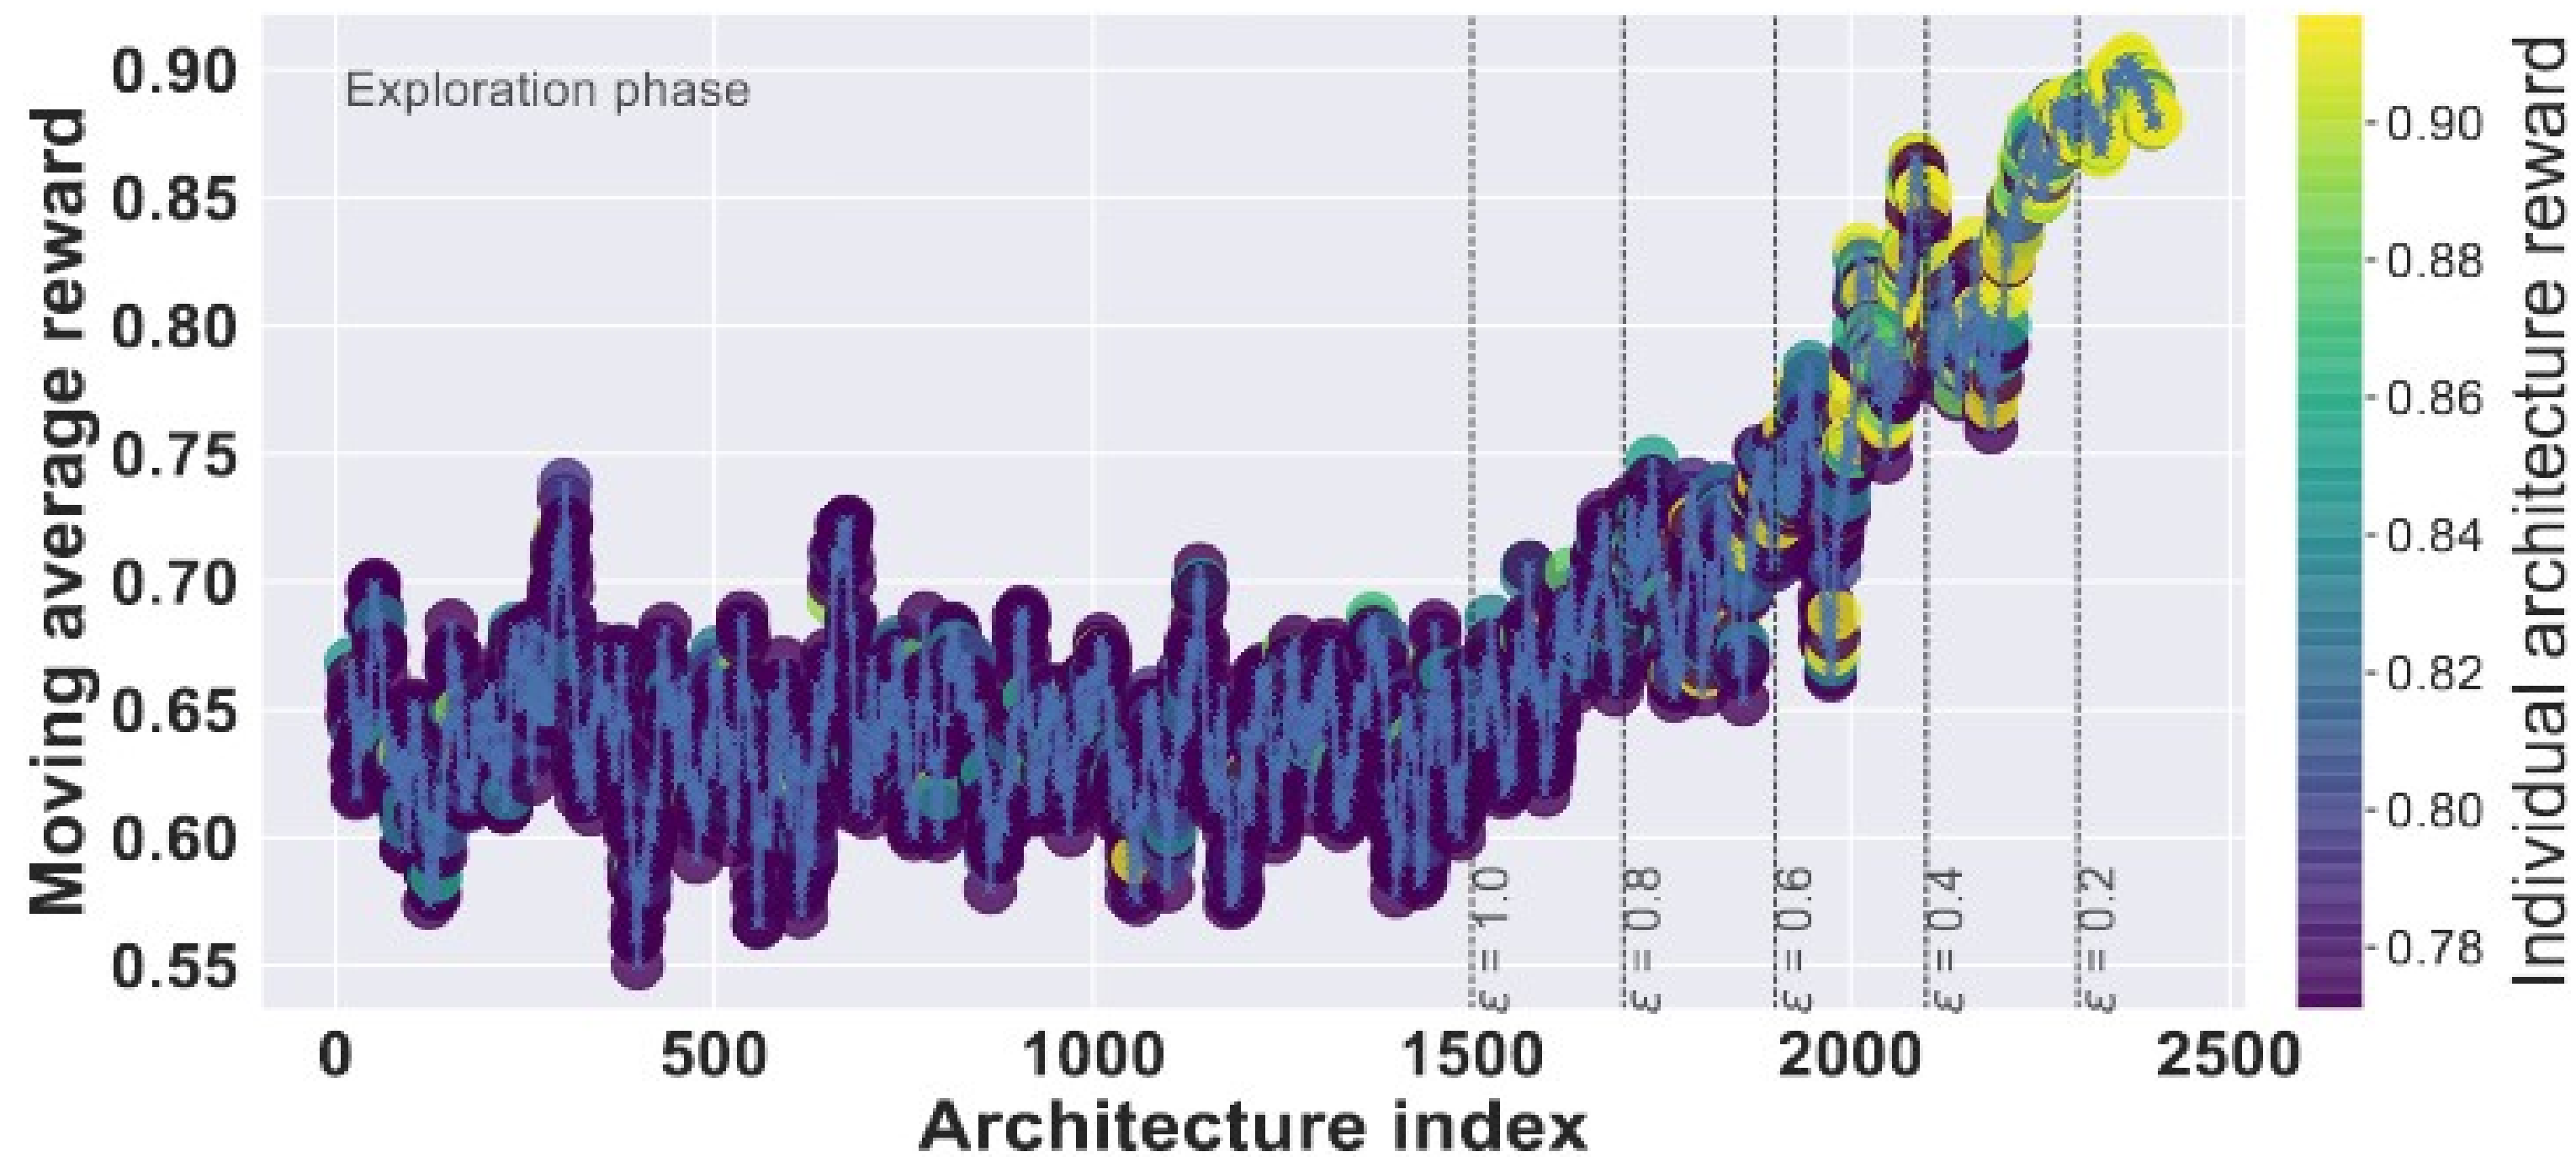
Meta-Learning Application: NAS of deep priors



Martin Mundt°, Iuliia Pliushch°, Visvanathan Ramesh (° equal contribution). Neural Architecture Search of Deep Priors: Towards Continual Learning without Catastrophic Interference. To appear in IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPR-W) 2021, Continual Learning in Computer Vision Workshop (CLVision)

Meta-Learning Application: NAS of deep priors


FashionMNIST DP-NAS



Martin Mundt[°], Iuliia Plushch[°], Visvanathan Ramesh ([°] equal contribution). Neural Architecture Search of Deep Priors: Towards Continual Learning without Catastrophic Interference. To appear in IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPR-W) 2021, Continual Learning in Computer Vision Workshop (CLVision)

Meta-Learning Application: NAS of deep priors

Random vs. trained FashionMNIST accuracy comparison



Martin Mundt°, Iuliia Pliushch°, Visvanathan Ramesh (° equal contribution). Neural Architecture Search of Deep Priors: Towards Continual Learning without Catastrophic Interference. To appear in IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPR-W) 2021, Continual Learning in Computer Vision Workshop (CLVision)